

CSC 330 / Spring 2023

OCaml
Patterns, Tail Recursion,
Exceptions

Ibrahim Numanagić /nu-ma-nagg-ich/

based on Dan Grossman's Lecture materials

Nested patterns

Nested patterns

- ▶ We can nest patterns inside other patterns
 - Just like we can nest expressions as deep as we want
 - Anywhere a variable can appear in our current patterns, we can put a pattern
 - Often avoids hard-to-read nested match expressions
- ▶ So full meaning of pattern-matching is to compare a pattern against a value for the *same shape* and bind variables to the *right parts*
 - More precise recursive definition coming after some examples

Code demonstration

Nested patterns

- ▶ Nested patterns are used in different contexts to express algorithms concisely and elegantly
 - Match on multiple things with same shape at same time
 - Go multiple levels into data structure at once
 - Use wildcard patterns when you do not need the data
 - Make match expressions that look like tables

Nested match expressions?

- Sometimes a match expression inside a match expression is a missed opportunity for nested pattern matching

```
match xs with
| [] -> 0
| x :: xs' -> match xs' with ...
```

- Other times it is necessary because you need to first compute with data extracted via the outer match

```
match xs with
| [] -> 0
| x :: xs' -> match f x s' with ...
```

- !: sometimes you need to wrap inner `match` statements with parentheses; otherwise you can end up with spectacularly confusing errors!

Patterns: generalization

► Syntax:

```
match e with p1 -> e1 | ... | pN -> eN  
let p1 = e in e1  
let [rec] f p1 = e1
```

where `e, e1, ..., eN` are expressions, and `p1, ..., pN` are patterns

► Pattern syntax: A pattern is one of:

- `x`: a variable
- `_`: a wildcard
- `(p1, ..., pN)`: a tuple of patterns
- `C p`: a constructor `C` applied to a pattern `p`

► Patterns are **NOT** expressions (though they kinda look like expressions)!

Match expressions: evaluation

Will now focus on match expressions, others are similar

- ▶ We need two things:
 1. Semantics of matching a pattern p to a value v
 - Does it, or does it not match?
 - If it matches, creates a set of (local) bindings
 2. Use (1) to define match expressions' semantics

Pattern-matching revealed

Yet another elegant recursive definition:

- ▶ Given pattern p and value v :
 - if p is a variable x , the match succeeds and x is bound to v ;
 - if p is $\underline{}$, the match succeeds and no bindings are introduced;
 - if p is (p_1, \dots, p_N) and v is (v_1, \dots, v_N) , the match succeeds if and only if p_1 matches v_1 , and ..., and p_N matches v_N (the bindings are the union of all bindings from the submatches);
 - if p is $C\ p_1$, the match succeeds if and only if v is $C\ v_1$ (i.e., the same constructor) and p_1 matches v_1 (the bindings are the bindings from the submatch)

Examples

- ▶ `a :: b :: c :: d` matches all lists with ≥ 3 elements
- ▶ `a :: b :: c :: []` matches all lists with exactly 3 elements
- ▶ `((a, b), (c, d)) :: e` matches all non-empty lists of pairs of pairs

Match expressions: semantics

► Evaluation:

```
match e with p1 -> e1 | ... | pN -> eN
```

1. Evaluate e to v
2. Check in p_1, p_2, \dots order; find first p_i that matches
3. Evaluate e_i in environment extended with bindings from the match
4. Result of (3) is the overall result

Note: This is the *semantics*, but the implementation can *optimize* with concepts like binary search

Match expressions: type checking

► Type checking:

```
match e with p1 -> e1 | ... | pN -> eN
```

- Type-check `e` to some type `t`
- Each `pi` must match against some value of type `t`
- Similar pattern, type matching definition to get the types of bindings
 - Typecheck `ei` in a static environment that has those bindings
- All `e1, ..., eN` must have the same type `t2`, which is the overall type
- Also, some fantastic features:
 - Error if a `pi` can never match because of earlier patterns (dead code)
 - Warning if a value of type `t` could match no `pi` (incomplete match)
 - (With nested patterns, these require some fairly fancy algorithms)

Code demonstration

Tail Recursion

Recursion

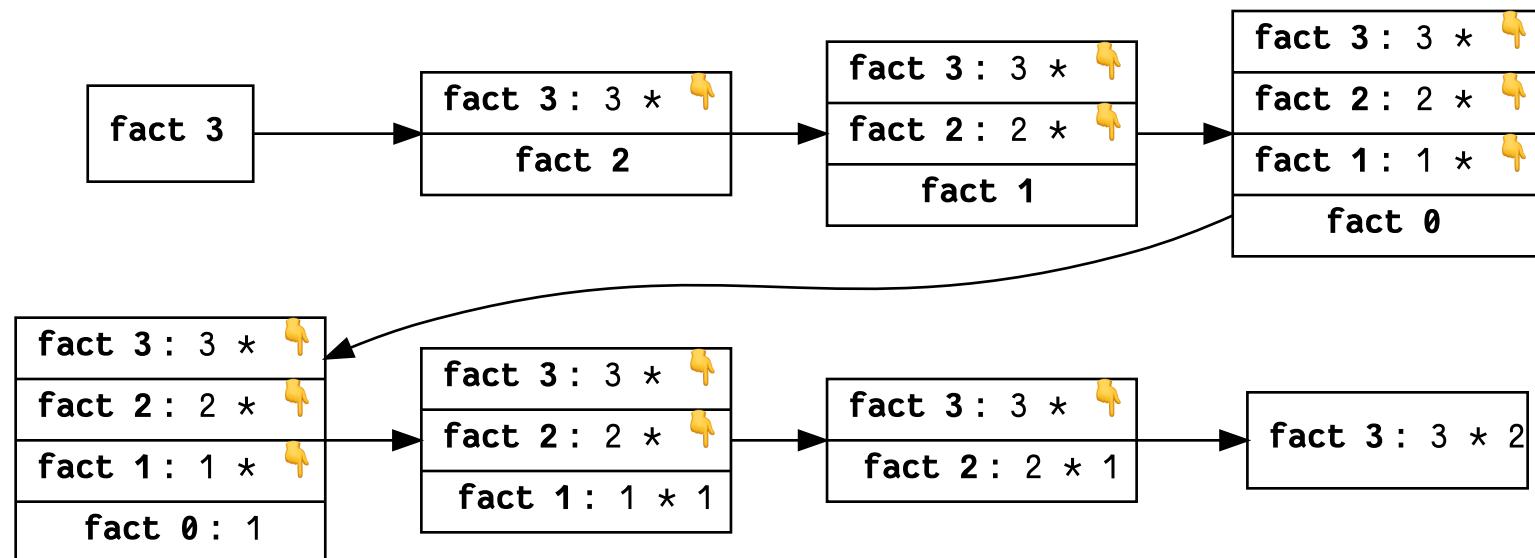
- ▶ Now have lots of practice with recursion
 - *Code follows the data*: recursive data processed recursively!
- ▶ No need for loops and assignment statements. Recursion often easier:
 - e.g., processing trees, appending lists
- ▶ But every recursive call needs to get “fresh space” for arguments and local variables, separate from the caller’s “space”
- ▶ **Next:** let’s refine our mental model of how calls execute (call stacks)
 - And see how recursion can be *as efficient as a loop!*

Call stacks

- ▶ While a program runs, there is a **call stack** for all function calls that have started but not yet completed
 - Calling `f` pushes a stack frame for `f` on the stack
 - When a call finishes, its stack frame is popped from the stack
 - Recursion: Can have more than one stack frames be for calls to the same function
 - ▶ A stack frame stores information: values of parameters, values of local variables, and “what’s left to do” for the function body
- ▶ Function calls finish in the reverse order that they start!

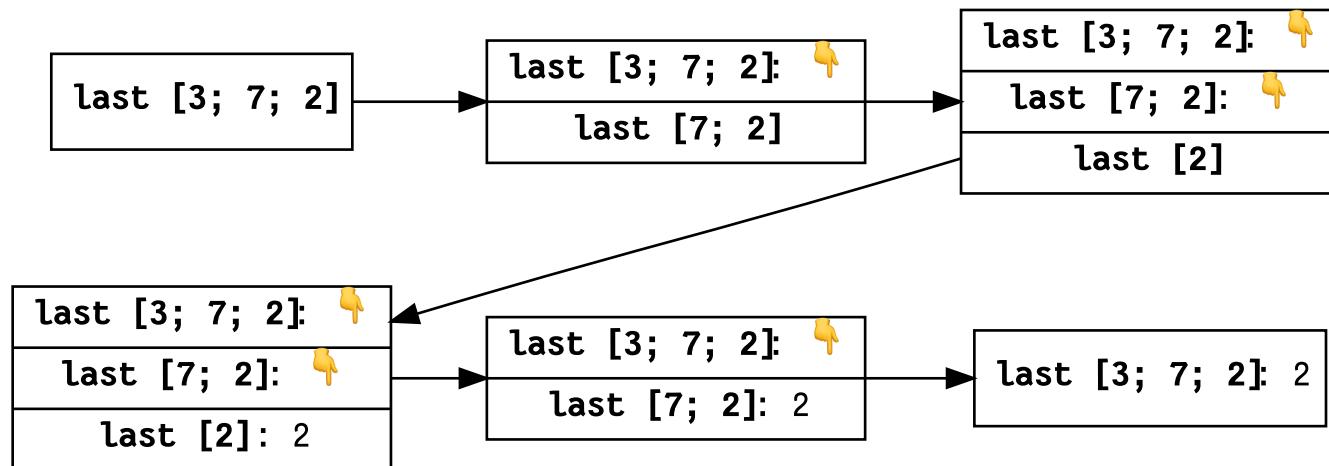
Example

```
let rec fact n =
  if n = 0 then 1 else n * fact(n - 1)
let x = fact 3
```



Example

```
let rec last xs =
  match xs with
  | x :: [] -> x
  | _ :: xs' -> last xs'
let x = last [3; 7; 2]
```



Do we *always* need a stack frame?

- ▶ **No!** It is unnecessary to keep around a stack-frame just so it can get a result from a function call and immediately return it “as is”
 - A call that will be the caller’s answer *as is* is called a **tail call**
 - A recursive function where all recursive calls are tail calls is called **tail recursive**
- ▶ OCaml handles tail calls specially: **tail-call optimization**
 - Pops caller’s frame **before** the call
 - This can take $O(n)$ space usage down to $O(1)$ 🎉
 - No *stack overflow* issues anymore!
 - Along with other optimizations, as efficient as a loop

last is tail-recursive!

```
let rec last xs =
  match xs with
  | x :: [] -> x
  | _ :: xs' -> last xs'
let x = last [3; 7; 2]
```



Making functions tail-recursive

- Sometimes we can rewrite our algorithm to be tail-recursive:

```
let fact n =
  let rec loop (n, acc) =
    if n = 0 then acc else loop (n - 1, acc * n)
  in
  loop(n, 1)
let x = fact 3
```

- Sometimes we can't easily or profitably do so:

- processing a tree (two recursive calls can't both be tail calls)
- if order matters (e.g., concatenating a string list)

Methodology

- ▶ There is a *methodology* to guide the transformation to a tail-recursive function:
 - Create a helper function that takes an **accumulator**
 - Old base case becomes initial accumulator
 - New base case becomes final accumulator
- ▶ Up to you to reason about whether this works
 - Example: for factorial, it is crucial that *multiplication is associative*

Code demonstration

Examples

- ▶ **Summing a list:** $O(1)$ stack space vs. length-of-list stack space
 - But notice both versions of factorial and list-sum take *time* linear in length of list
- ▶ **Reversing a list:** Tail-recursive version actually *much better*
 - Non-tail recursive version is *quadratic time* $O(n^2)$:
 - ▶ each recursive call uses `append`, which must traverse the list;
 - ▶ so total time for the appends is $1 + 2 + \dots + (n + 1)$ which is roughly $\frac{n^2}{2}$
 - Moral: beware appending to lists, especially within outer recursion
 - Consing is constant time and fast (why?), so tail-recursive version is linear time

Moral: don't overdo it!

- ▶ Tail-recursion is not always feasible or elegant
- ▶ Tail-recursion is not always necessary: beware premature optimization and favor readability
- ▶ But where it is feasible and efficiency matters, it is a key tool for functional programmers to get the efficiency of loops
 - If you can write it as a loop, you can write it as a tail-recursive function: *updated values* are arguments to the tail call!
- ▶ Most functional languages, including OCaml, promise tail-call optimization

Precise definition of tail calls

- ▶ How do you know if a call is a tail call?
 - Humans: “Oh, I see the caller has no more work to do after it gets the result.”
 - Semanticists and compilers: a tail call is a function call that *is in tail position*
 - ▶ which is defined recursively over the syntax of expressions...

Tail position

- ▶ In `let f p = e`, the expression `e` is in tail position
- ▶ If `if e1 then e2 else e3` is in tail position:
 - then: `e2` and `e3` are in tail position;
 - ▶ but `e1` is **not** in tail position!
- ▶ If `let p = e1 in e2` is in tail position:
 - then: `e2` is in tail position
 - ▶ but `e1` is **not** in tail position
- ▶ If `e1 + e2`, the subexpressions are **not** in tail position
- ▶ ... (more cases for every kind of expression)
- ▶ If `e` is **not** in tail position, then neither is any of its subexpressions

Exceptions

Exceptions

- ▶ Exception bindings declare new kinds of exception:

```
exception Whatever  
exception BadNum of int
```

- ▶ `raise` expressions can throw an exception:

```
raise Whatever  
raise (BadNum 1)
```

- ▶ `try ... with ...` expressions can catch exceptions:

```
try e with Whatever -> 0  
try e with BadNum n -> n
```

Exceptions

- ▶ Exception bindings declare new kinds of exception:

```
exception Whatever  
exception BadNum of int
```

- ▶ We *build* exceptions, which are just *values* (!) with these constructors:

```
Whatever  
(BadNum 1)
```

- ▶ That is different than *raising* an exception with the `raise` expression:

```
raise Whatever  
raise (BadNum 1)
```

A bit more precisely

- ▶ New exception bindings add new variants to “the one exception type” `exn`
- ▶ We build values of type `exn` just like we build variant types with constructors
 - Can pass them around, though that isn’t that common
- ▶ We raise (throw) exceptions with `raise e`
 - **Type checking:**
 - ▶ `e` must have type `exn`
 - ▶ Result type is *any type you want* (!)
 - **Evaluation:**
 - ▶ Do not produce a result (!); but...
 - ▶ *cancel everything and pass the exception produced by `e` to the nearest try expression on the call stack*
 - ▶ and keep doing this while you can!

Try expressions

- ▶ `try ... with ...` expressions can *catch* exceptions:

```
try e1 with Whatever -> e2  
try e1 with BadNum n -> e2
```

- ▶ Evaluation:

- Just evaluate `e1` and that is the result
 - But if `e1` raises an exception *and* that exception *matches* the pattern, then evaluate `e2` and that is the result

- ▶ Type checking:

- `e1` and `e2` must have the same type and that is the overall type

Code demonstration

